



Databases and JDBC

J. Pereira

jop@di.uminho.pt

R. Vilaça


rmv@di.uminho.pt



Universidade do Minho

Motivation

“select * from t where a = 0”



```
Connection c = new PostgreSQL(...);  
List<Row> r = c.execute("select * from ...");
```

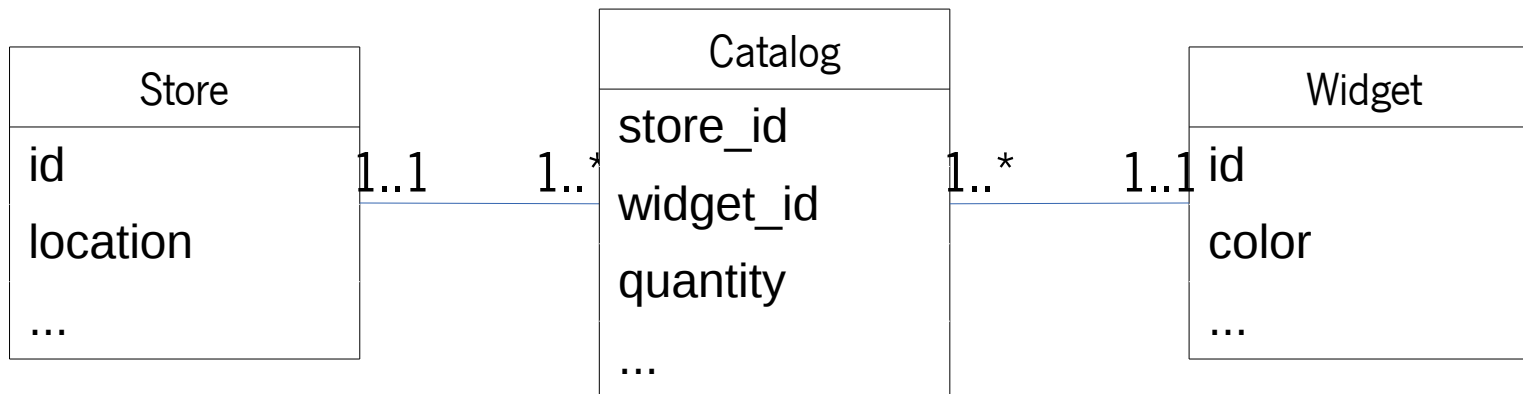
Why not?

Overview

- › Example
- › Query optimization and compilation
- › Query execution
- › Transactions and isolation
- › More...
 - › Distributed databases research @ HASLab
 - › Distributed databases @ MEI

Example

- › The colorful widget finder:
 - › Suppliers at multiple locations
 - › Widgets come in different colors
- › Operations:
 - › “Buy a widget.”
 - › “Where to find red widgets in Porto?”



Example

> “Buy widget.”

<i>id</i>	<i>location</i>
1	Rome
2	Lisbon
3	Porto
4	Lisbon
5	Brussels
6	Lisbon
7	Lisbon
8	Porto
...	...

<i>s_id</i>	<i>w_id</i>	<i>quan</i>
1	4	...
2	3	...
2	1	...

<i>id</i>	<i>color</i>
1	green
2	red
3	green

```
SELECT quantity FROM Catalog
    WHERE s_id = ... AND w_id = ...
if (q > 0)
    UPDATE Catalog
        SET quantity = quantity - 1
        WHERE s_id = ... AND w_id = ...
```

Example

› “Buy widget.”

<i>id</i>	<i>location</i>
1	Rome
2	Lisbon
3	Porto
4	Lisbon
5	Brussels
6	Lisbon
7	Lisbon
8	Porto
...	...

<i>s_id</i>	<i>w_id</i>	<i>quan</i>
1	4	...
2	3	...
2	1	...
2	2	...
3	4	...
4	4	...
...

<i>id</i>	<i>color</i>
1	green
2	red
3	green
4	red
5	blue
...	...

```
UPDATE Catalog
  SET quantity = quantity - 1
  WHERE s_id = ... AND w_id = ...
        AND quantity > 0
```

Example

› “Where to find red widgets in Porto?”

<i>id</i>	<i>location</i>
1	Rome
2	Lisbon
3	Porto
4	Lisbon
5	Brussels
6	Lisbon
7	Lisbon
8	Porto
...	...

<i>s_id</i>	<i>w_id</i>	<i>quan</i>
1	4	...
2	3	...
2	1	...
2	2	...
3	4	...

<i>id</i>	<i>color</i>
1	green
2	red
3	green
4	red
5	...

```
SELECT * FROM Store WHERE location='Porto'
for(...)
    SELECT * FROM Catalog WHERE s_id = ...
for(...)
    SELECT * FROM Widget WHERE w.color='red'
```

Example

> “Where to find red widgets in Porto?”

<i>id</i>	<i>location</i>
1	Rom
2	Lisbon
3	Porto
4	Lisbon

<i>s_id</i>	<i>w_id</i>	<i>quan</i>
1	4	...
2	3	...
2	1	...
2	2	...


<i>id</i>	<i>color</i>
1	green
2	red
3	green
4	red
5	blue
...	...

```
SELECT
    s.id, s.location
FROM
    Store s
    INNER JOIN Catalog c ON s.id=c.sid
    INNER JOIN Widget w ON c.wid=w.id
WHERE
    s.location='Porto' AND w.color='red'
```


Query processing

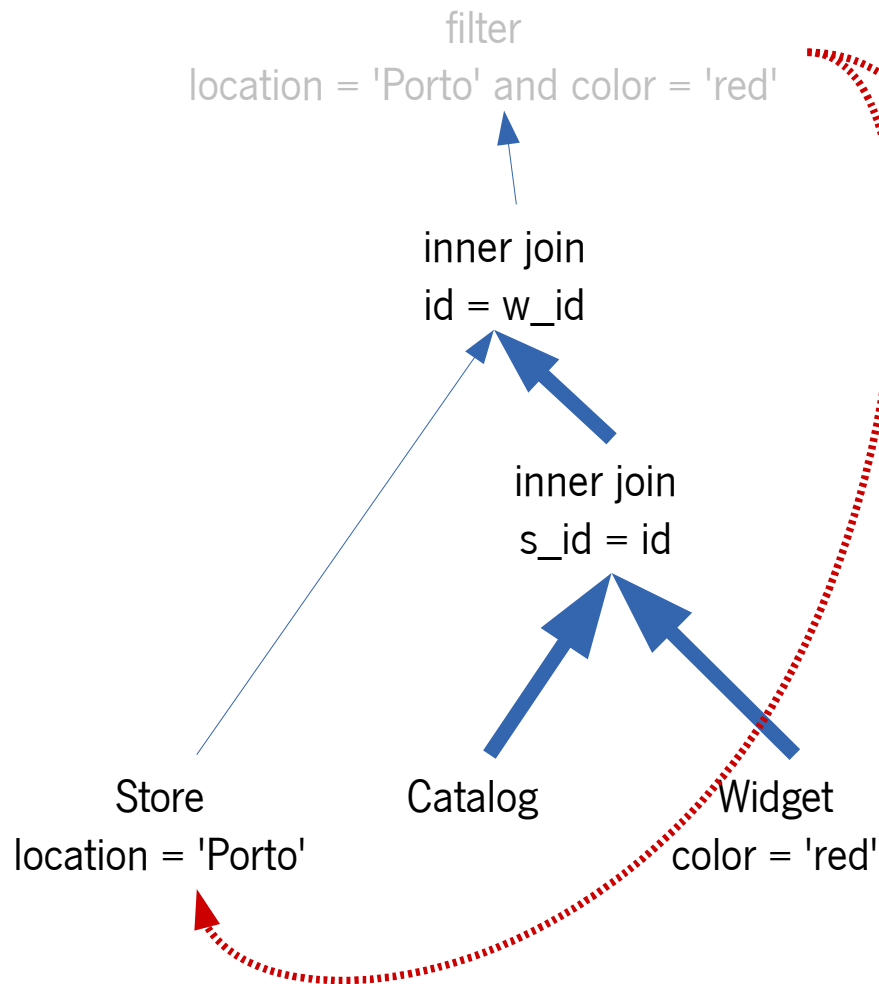
- › Query compilation to a relational operator expression:

```
SELECT
    s.id, s.location
FROM
    Store s
    INNER JOIN Catalog c ON s.id=c.s_id
    INNER JOIN Widget w ON c.w_id=w.id
WHERE
    s.location='Porto' AND w.color='red'
```




Query processing

› Heuristic optimization:




Query processing

- › Cost-based optimization using:
 - › Estimated row counts
 - › Row cost



VS.




Conclusions

- › SQL is declarative
 - › Describes “What to do”, not “How to do it”
 - › The DBMS decides “How to do it”
 - › The DBMS often knows best!
- › Consequences:
 - › Let the DBMS do as much as possible (prefer slide 8 to 7)
 - › Use “prepare” judiciously
 - › Meta-data is available before execution

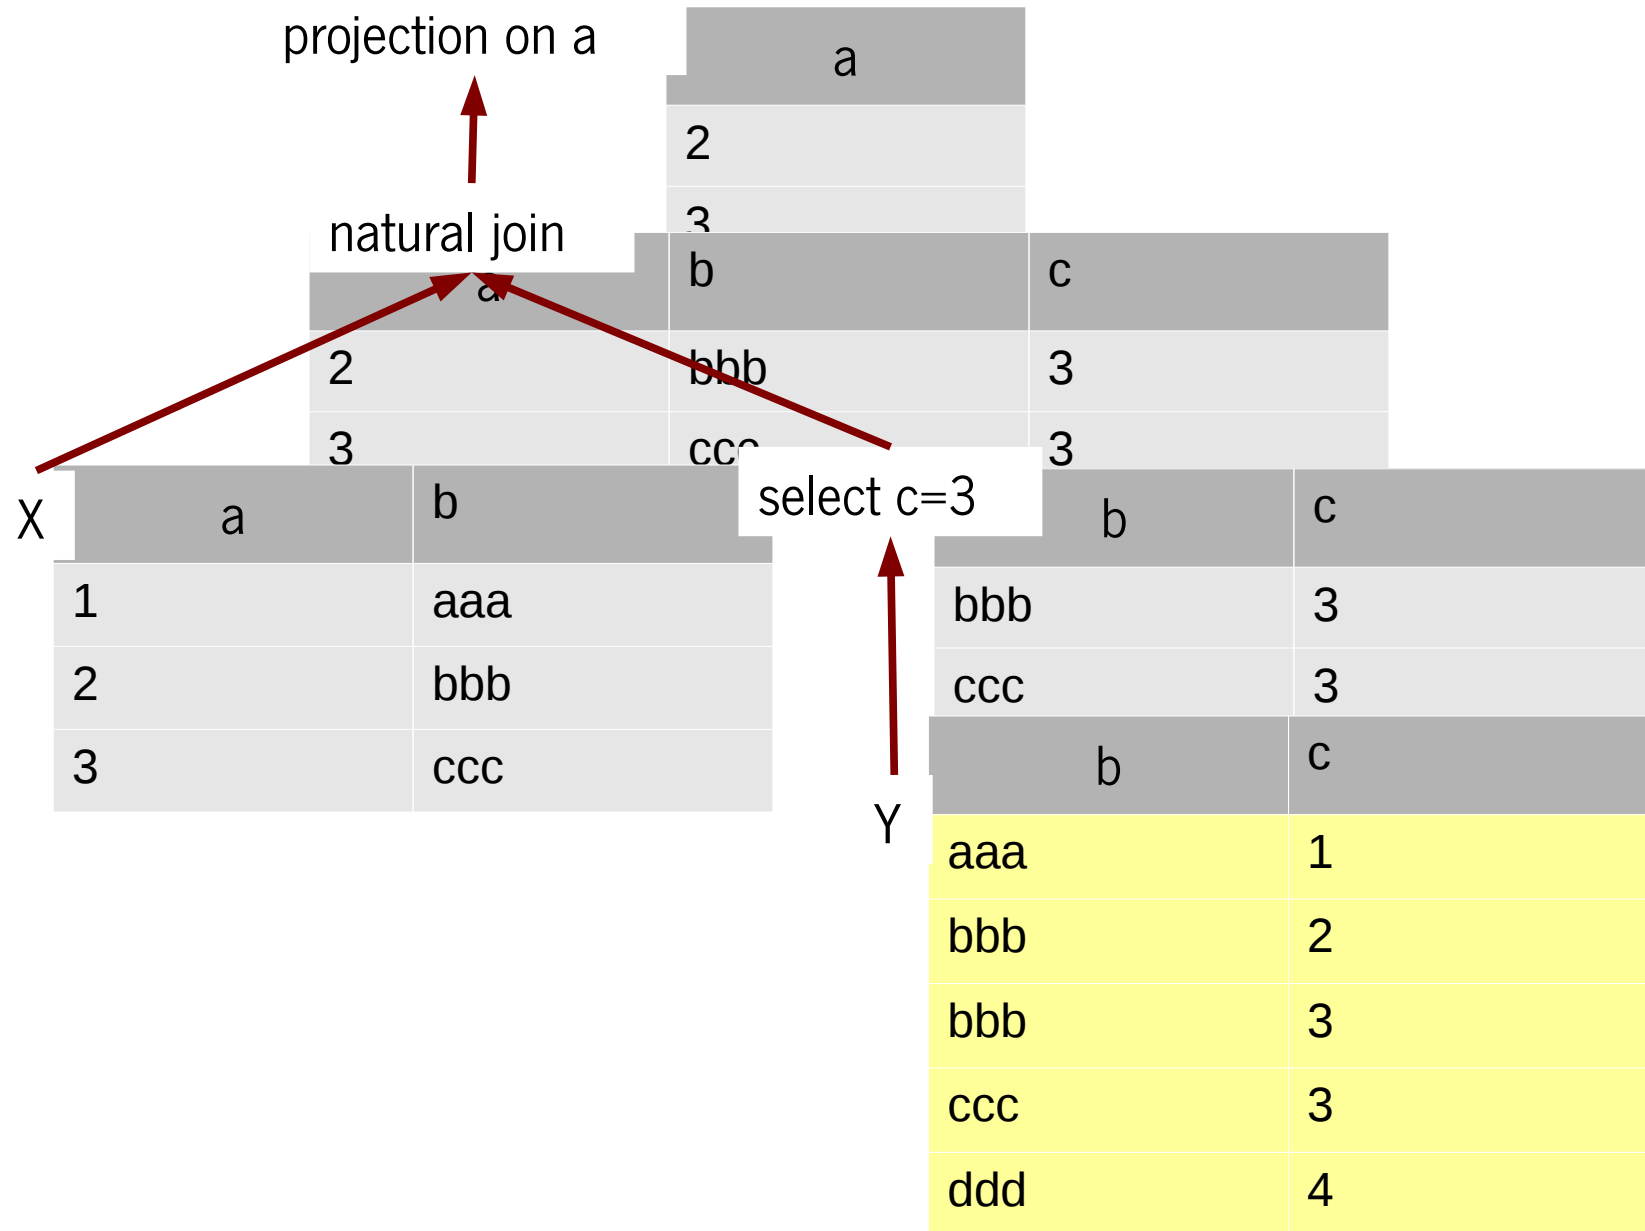
Query execution

- › Each operator is a function:
 - › Returns a relation
 - › Parameters are other relations
- › Computation order:
 - › From leaves to root



SELECT a FROM x NATURAL JOIN y WHERE c = 3

Execution with materialization



Execution with materialization

		a	
		2	
		3	
a		b	c
2		bbb	3
3		ccc	3
a	b		
1	aaa		
2	bbb		
3	ccc		
		b	c
		bbb	3
		ccc	3
		b	c
		aaa	1
		bbb	2
		bbb	3
		ccc	3
		ddd	4

Execution with materialization

The diagram illustrates the decomposition of a 3D tensor into two 3D tensors and a 3D tensor, showing the resulting 3D tensors and their dimensions.

Top 3D Tensor (Decomposition):

- Dimensions: 2, 3, 3
- Structure: A 3D tensor with dimensions 2, 3, and 3. The first dimension is 2, the second is 3, and the third is 3. The tensor is divided into two parts: a 2x3x3 part and a 2x3x3 part.

Bottom 3D Tensor (Decomposition):

- Dimensions: 2, 3, 3
- Structure: A 3D tensor with dimensions 2, 3, and 3. The first dimension is 2, the second is 3, and the third is 3. The tensor is divided into two parts: a 2x3x3 part and a 2x3x3 part.

Resulting 3D Tensors:

- Left 3D Tensor:** Dimensions 2, 3, 3. Structure: A 3D tensor with dimensions 2, 3, and 3. The first dimension is 2, the second is 3, and the third is 3. The tensor is divided into two parts: a 2x3x3 part and a 2x3x3 part.
- Right 3D Tensor:** Dimensions 2, 3, 3. Structure: A 3D tensor with dimensions 2, 3, and 3. The first dimension is 2, the second is 3, and the third is 3. The tensor is divided into two parts: a 2x3x3 part and a 2x3x3 part.

Execution with materialization

		a	
		2	
		3	
a		b	c
2		bbb	3
3		ccc	3
a	b		
1	aaa		
2	bbb		
3	ccc		
		b	c
		bbb	3
		ccc	3
		b	c
		aaa	1
		bbb	2
		bbb	3
		ccc	3
		ddd	4

Execution with materialization

		a			
		2			
		3			
	a	b	c		
	2	bbb	3		
	3	ccc	3		
a	b		b	c	
1	aaa		bbb	3	
2	bbb		ccc	3	
3	ccc		b	c	
			aaa	1	
			bbb	2	
			bbb	3	
			ccc	3	
			ddd	4	

Execution with materialization

		a	
		2	
		3	
a		b	c
2		bbb	3
3		ccc	3
a	b		
1	aaa		
2	bbb		
3	ccc		
		b	c
		bbb	3
		ccc	3
		b	c
		aaa	1
		bbb	2
		bbb	3
		ccc	3
		ddd	4

Execution with materialization

		a			
		2			
		3			
		a	b	c	
2		bbb		3	
3		ccc		3	
a		b		b	
1	aaa		bbb		3
2	bbb		ccc		3
3	ccc		b		c
		aaa		1	
		bbb		2	
		bbb		3	
		ccc		3	
		ddd		4	

Iteration

- › Each operator is an iterator
 - › Typical interface:
 - › `has_next()` / `next()`
- › Computation order:
 - › From leaves to root, for each record

Execution with iteration

		a			
		2			
		3			
a		b		c	
2		bbb		3	
3		ccc		3	
a	b			b	c
1	aaa			bbb	3
2	bbb			ccc	3
3	ccc			b	c
				aaa	1
				bbb	2
				bbb	3
				ccc	3
				ddd	4

Execution with iteration

		a			
		2			
		3			
a		b		c	
2		bbb		3	
3		ccc		3	
a	b			b	c
1	aaa			bbb	3
2	bbb			ccc	3
3	ccc			b	c
				aaa	1
				bbb	2
				bbb	3
				ccc	3
				ddd	4


Execution with iteration

		a	
		2	
		3	
a		b	c
2		bbb	3
3		ccc	3
a	b		
1	aaa		
2	bbb		
3	ccc		
		b	c
		bbb	3
		ccc	3
		b	c
		aaa	1
		bbb	2
		bbb	3
		ccc	3
		ddd	4

Execution with iteration

		a			
		2			
		3			
a		b		c	
2		bbb		3	
3		ccc		3	
a	b				
1	aaa				
2	bbb				
3	ccc				
		b		c	
		bbb		3	
		ccc		3	
		b		c	
		aaa		1	
		bbb		2	
		bbb		3	
		ccc		3	
		ddd		4	

Execution with iteration



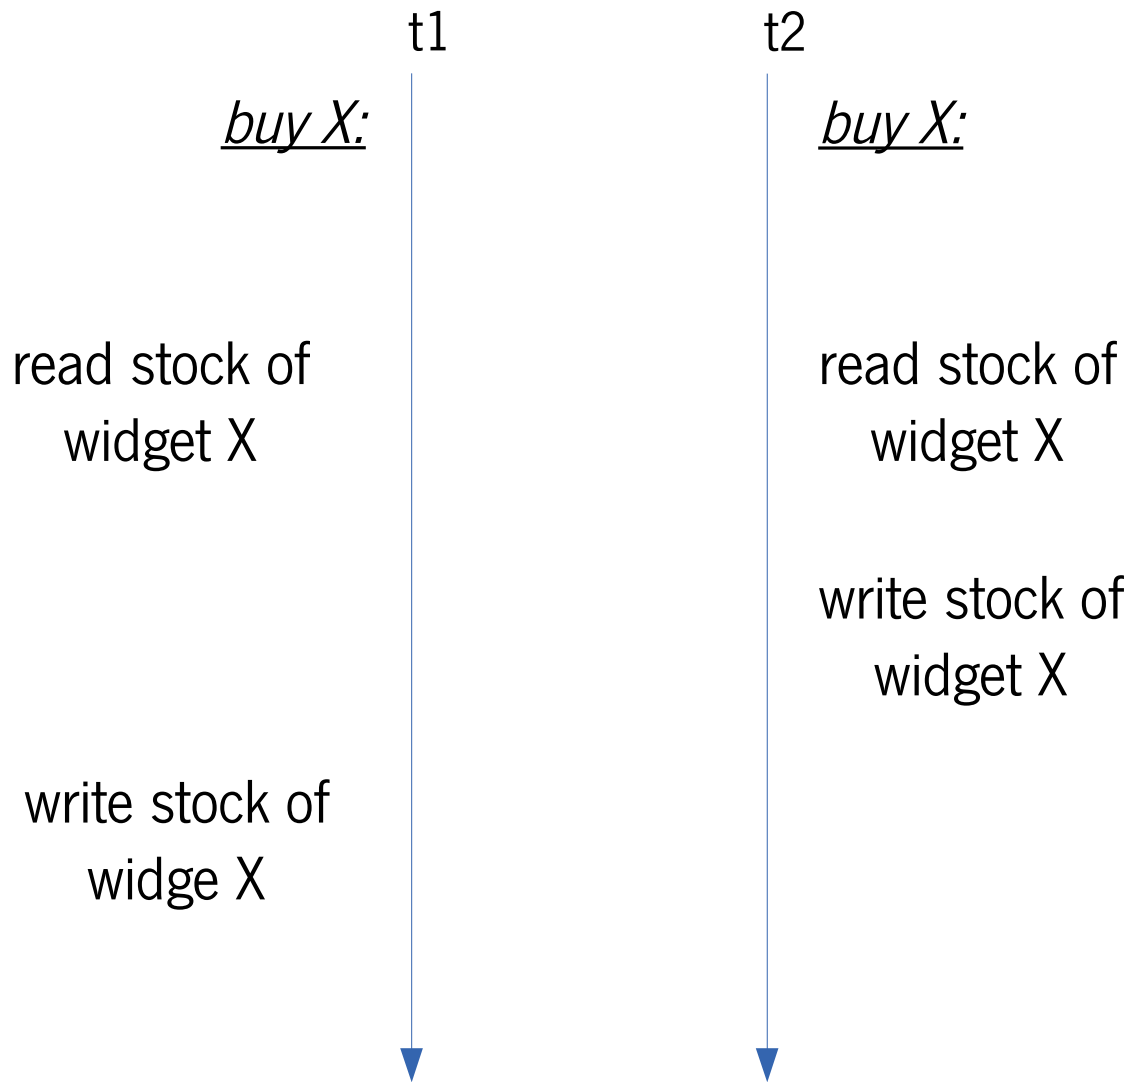
Execution with iteration

		a			
		2			
		3			
a		b		c	
2		bbb		3	
3		ccc		3	
a	b			b	c
1	aaa			bbb	3
2	bbb			ccc	3
3	ccc			b	c
				aaa	1
				bbb	2
				bbb	3
				ccc	3
				ddd	4

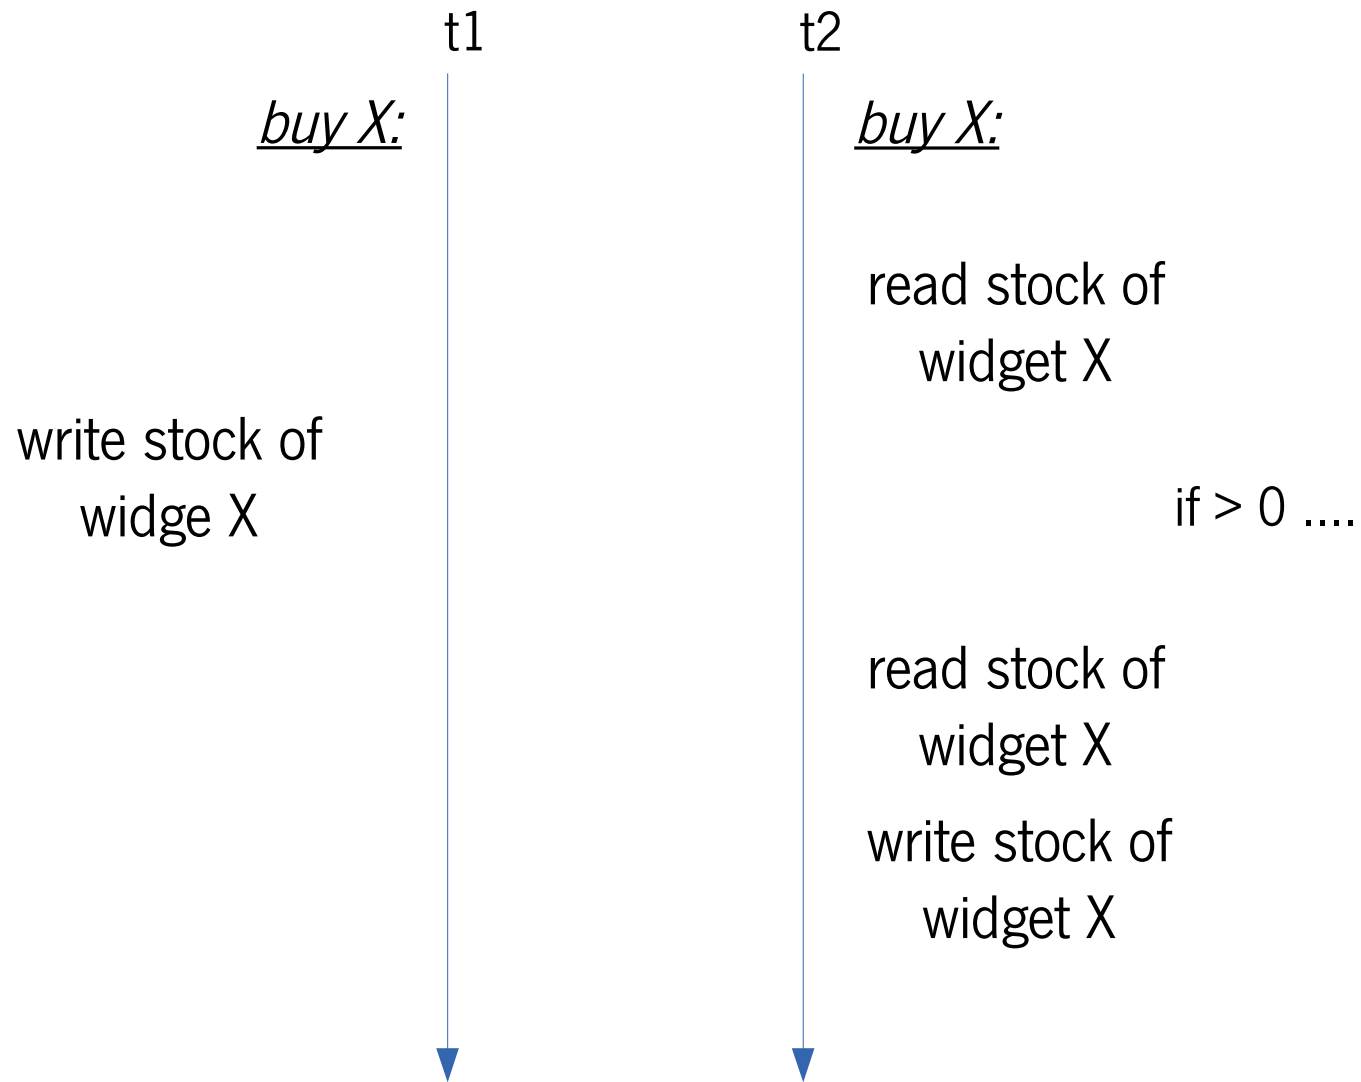
Conclusions

- › The DBMS is lazy:
 - › Query execution happens on demand
- › Consequences:
 - › At most one active ResultSet in the same Connection
 - › Don't force it to work when not needed
 - › Avoid operations that cannot be executed on demand
 - › Sort
 - › Aggregates

Lost update




Non-repeatable read




General problem

- › Transaction execution overlaps in time
- › In detail, some operation C should not be happening between operations A and B...




General approach

- › Postpone C
 - › t1 precedes t2




General approach

- › Remove C and related operations
 - › t1 executes alone



General approach

- › Anticipate C (i.e. execute C before the application has requested it!)
 - › t2 precedes t1
- › (Can you propose a mechanism to do this!?)



Conclusions

- › The DBMS deals with concurrent data access and modification
- › Consequences:
 - › Delimit transactions and set isolation level
 - › Expect (unexpected) rollbacks
 - › Take advantage of rollback to simplify code (prefer slide 6 to 5)



Database research @ HASLab




- › CumuloNimbo: Highly Scalable Transactional Multi-Tier PaaS

- › <http://cumulonimbo.eu>

- › CoherentPaaS: A Coherent and Rich PaaS with a Common Programming Model

- › <http://coherentpaas.eu>



CoherentPaaS

- › LeanBigData: Ultra-Scalable and Ultra-Efficient Integrated and Visual Big Data Analytics

- › <http://leanbigdata.eu>



<https://www.youtube.com/watch?v=2ShTc0v-3Aw>


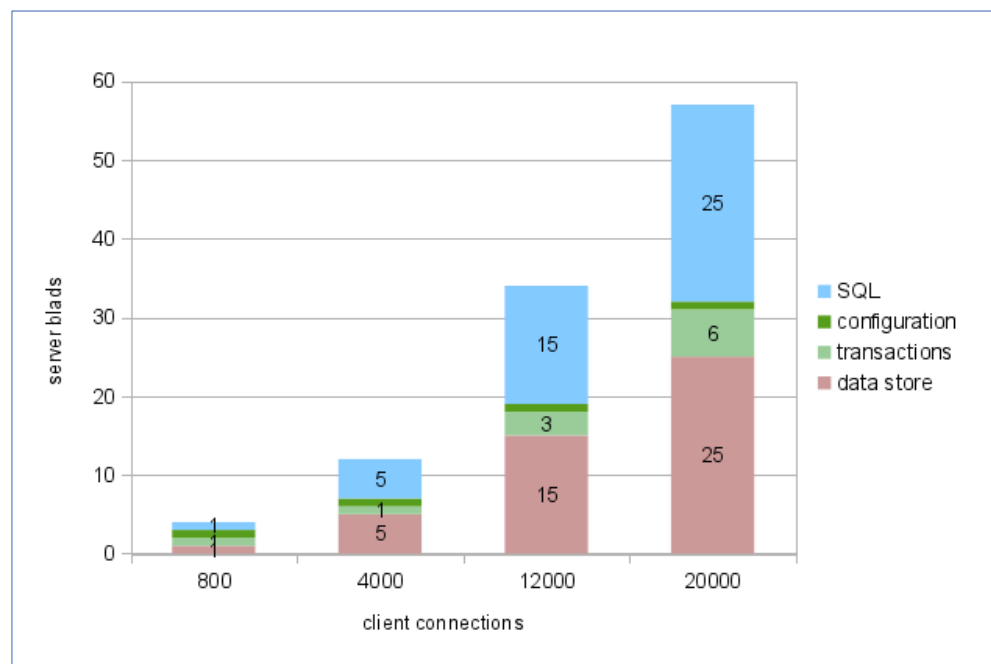
<https://www.youtube.com/watch?v=Z8l1Y57CCyQ>

<https://www.youtube.com/watch?v=R9QTmjxAdOk>



Database research @ HASLab

- › Standard SQL / JDBC
- › Independently elastic layers
- › Evaluated with TPC-C (OLTP standard)



quad-core Intel Xeon X3220 @ 2.40GHz / 8GB RAM
500GB SATA disk / 1Gbits Ethernet

Distributed databases @ MEI/MIEI

- › Distributed Systems and Cryptography
 - › Transactional Distributed Systems
 - › Fault Tolerance and Large Scale Distributed Systems
- › Applications Engineering
 - › Database Administration



HIGH-ASSURANCE
SOFTWARE LABORATORY

IMPROVING PRACTICE THROUGH THEORY