

# Query Processing and JDBC

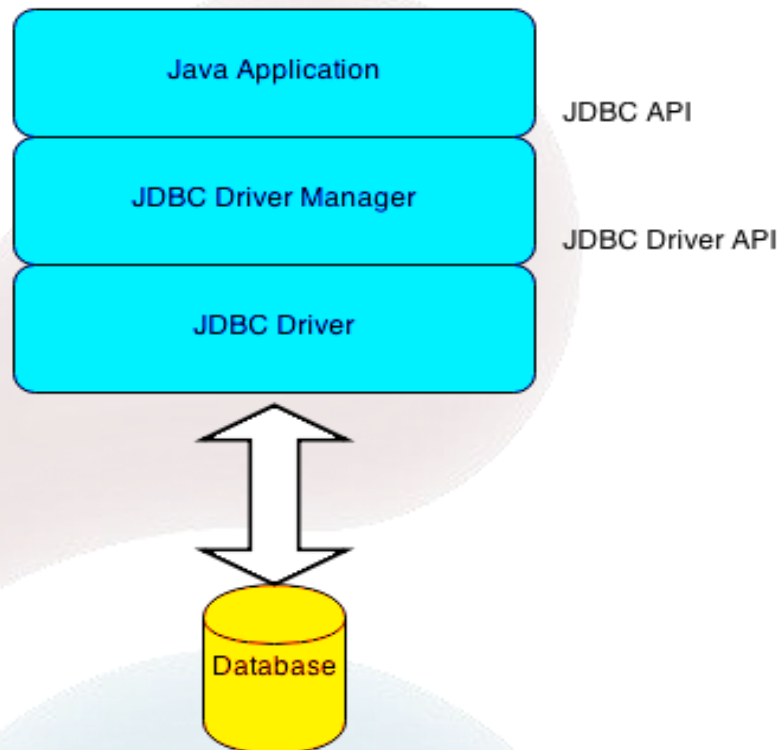
J. Pereira

R. Vilaça

B. Sonntag

# JDBC

- › JDBC - Java Database Connectivity
- › Set of Java interfaces and classes in package `java.sql` to access relational databases.
- › SQL code is used implicitly inside Java code as Strings.



# JDBC – Initialization

## › Initialization:

- › Load and register JDBC driver in Driver Manager;
- › Obtain a Connection to the database using an URL;
- › It's also possible to use a DataSource instead.

## › Connections are:

- › Single-threaded;
- › Heavy weight;
- › Should be closed when no longer needed.

# JDBC - Drivers

- › The Driver interface is used by the vendors that implement JDBC drivers.
- › It is needed to load the driver class in the application. Can be done using `Class.forName("org.postgresql.Driver");`
- › The DriverManager class handles objects of type Driver
  - › Register, remove and list drivers
- › Use the DriverManager to obtain a Connection:
  - › `Connection connection = DriverManager.getConnection(url, user, pw);`

# JDBC – URLs

- › Uses syntax similar to net URLs

  - › jdbc:mysql://localhost:3306/mydb

  - › jdbc:derby:test//localhost/testdb

- › General syntax

  - › jdbc:subprotocol\_name:other\_stuff

  - › Where subprotocol selects the specific driver

  - › Format for other\_stuff depends on subprotocol, but in general:

    - › jdbc:subprotocol://hostname:port/database\_name

  - › jdbc:postgresql://localhost:5432/colorfulWidget would connect to the colorfulWidget database of localhost, using the Postgres driver

# JDBC - Statement

- › Used to execute SQL.
- › Created using the `createStatement()` method on a `Connection` object.
  - › `Statement stmt = connection.createStatement();`
- › Executing DDL statements: CREATE, DROP, ALTER TABLE
  - › `stmt.execute("CREATE TABLE widget (id INT PRIMARY KEY, color CHAR(20))");`
- › Executing DML update statements: INSERT, UPDATE
  - › `int modifiedRows = stmt.executeUpdate("INSERT INTO widget (id,color) VALUES (1, 'Red')");`
  - › Returns the number of modified rows.

# JDBC – Prepared Statements

- › Compiled once and executed several times. More efficient when several similar queries with different parameters are executed
- › Accepts SQL queries with parameters specified as “?”.
- › The parameters are set using the setXXX() where XXX is the parameter data type.
  - › `PreparedStatement stmt = connection.prepareStatement("INSERT INTO widget (id, color) VALUES (?, ?)");`
  - › `stmt.setInt(1, 1);`
  - › `stmt.setString(2, "Red");`

# JDBC - Queries

- › Queries returning data are executed using `executeQuery`:
  - › `ResultSet rs = stmt.executeQuery("SELECT * FROM widget");`
- › The `ResultSet` is similar to the `Iterator`:
  - › Allows to navigate through the retrieved data and read the columns selected in the query;
  - › Main methods: `next()`, `previous()`, `first()` and `last()`.
  - › Columns are fetched using the appropriate method for the type of that column and either the column number or the column name:
    - › `rs.getInt("id");`
- › Use Prepared Statements to avoid compilation overhead.



# JDBC - Isolation

- › Delimiting transactions:
  - › `setAutoCommit(false)` and `commit()/rollback()`
  - › Undo modifications on rollback
  - › Be prepared for conflicting transactions
- › The application must be ready for concurrency aborts
  - › Handling possible errors

# JDBC – SQL/Java Types

<i>SQL Type</i>	<i>Java Type</i>
CHAR	String
VARCHAR	String
LONGVARCHAR	String
NUMERIC	java.Math.BigDecimal
DECIMAL	java.Math.BigDecimal
BIT	boolean
TINYINT	int
SMALLINT	int
INTEGER	int
BIGINT	long
REAL	float
FLOAT	double
DOUBLE	double
BINARY	byte[]
VARBINARY	byte[]
DATE	java.sql.Date
TIME	java.sql.Time
TIMESTAMP	java.sql.Timestamp

# Distributed databases @ MEI

- › Distributed Systems and Cryptography
  - › Reliable Distributed Systems
- › Applications Engineering
  - › Database Administration

# Research @ HASLab

## › CumuloNimbo: Highly Scalable Transactional Multi-Tier PaaS

SQL Query Engine

Scalable  
Transactions

NoSQL  
Data Store


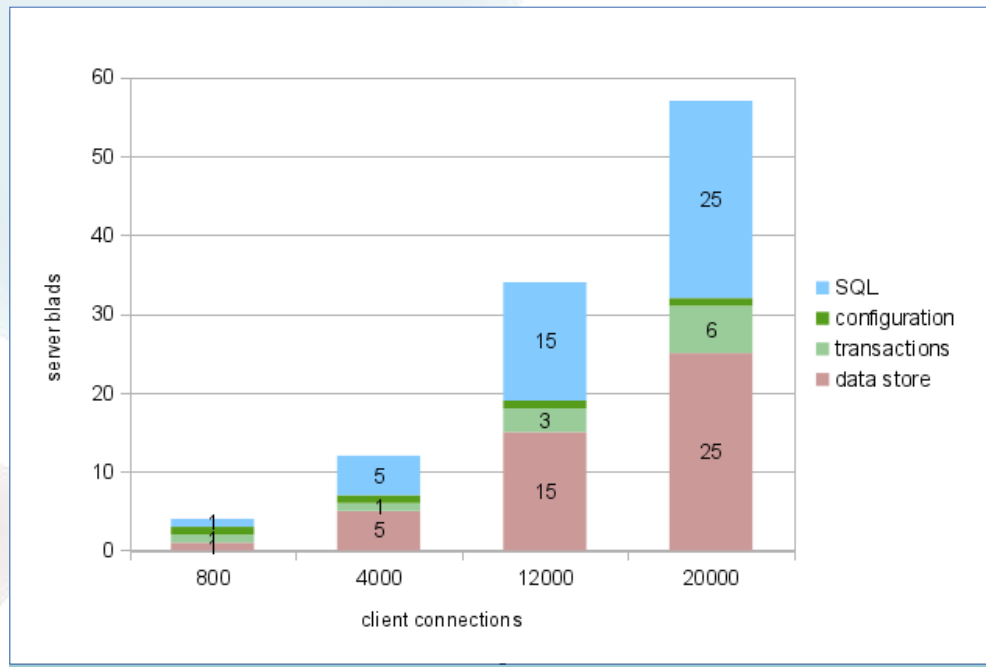


<https://www.youtube.com/watch?v=2ShTc0v-3Aw>

# Research @ HASLab

## › CumuloNimbo:

- › Standard SQL / JDBC
- › Independently elastic layers
- › Evaluated with TPC-C (OLTP standard)



quad-core Intel Xeon X3220 @ 2.40GHz / 8GB RAM  
500GB SATA disk / 1Gbits Ethernet

# Research @ HASLab



## › CoherentPaaS

- › A Coherent and Rich PaaS with a Common Programming Model

## › LeanBigData

- › Ultra-Scalable and Ultra-Efficient Integrated and Visual Big Data Analytics



HIGH-ASSURANCE  
SOFTWARE LABORATORY

# IMPROVING PRACTICE THROUGH THEORY